

Movie Recommendation System

**Project By
Tejas Shinde**

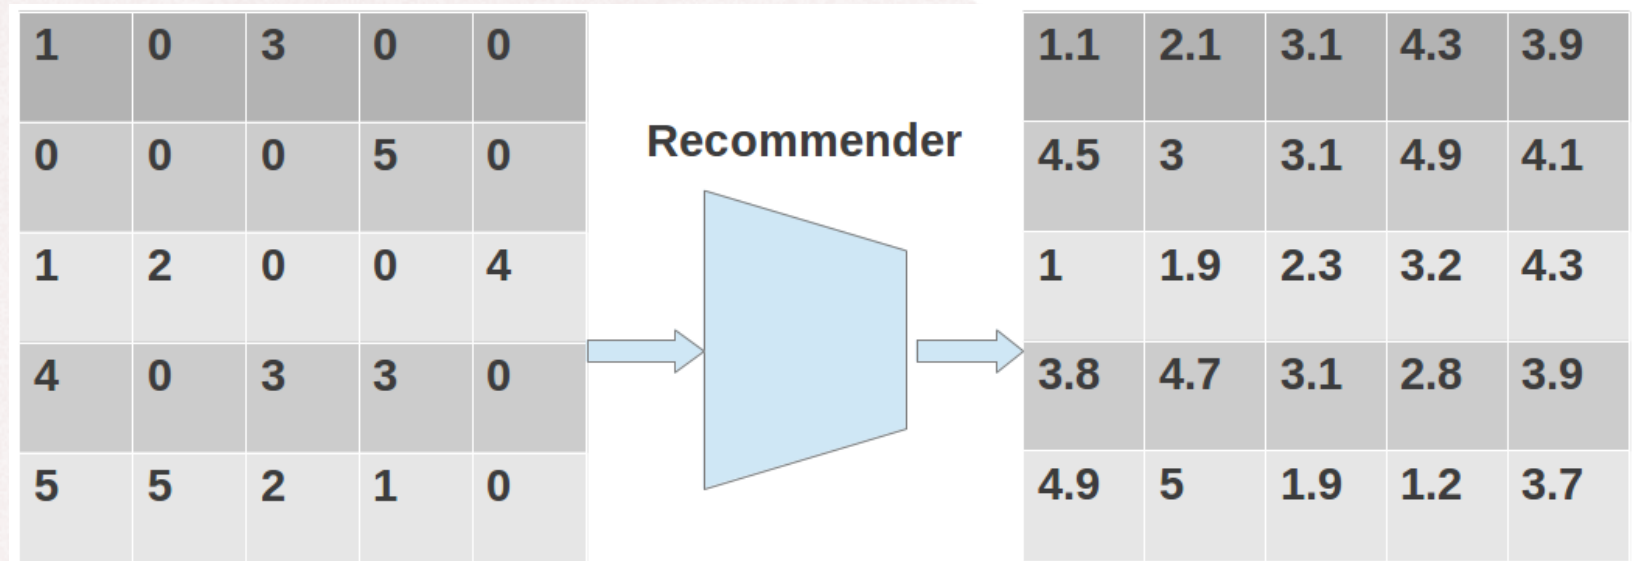
Problem Description

- Group of users
- Set of movies
- Every user has watched some movies but not all
- If users watch a movie, they rate it as well

Problem Statement

- Given a user how to recommend a new movie using the ratings?

Problem Formulation



Matrix M with N_u rows and N_m columns

Possible Approaches

- Collaborative Filtering
 - K-Nearest Neighbors
 - Using cosine similarity metric

$$\text{sim}(m_1, m_2) = \frac{\langle \mathbf{r}_{m1} - \mathbf{r}_{\text{mean}}, \mathbf{r}_{m2} - \mathbf{r}_{\text{mean}} \rangle}{\|\mathbf{r}_{m1} - \mathbf{r}_{\text{mean}}\|_2 \times \|\mathbf{r}_{m2} - \mathbf{r}_{\text{mean}}\|_2}$$

- Low Rank Approximation
 - $\mathbf{R} = \mathbf{U} \mathbf{V}^T$ for some lower dimensional feature space S_F .

Implementation

- 'n' users, 'm' movies
- discover the latent 'k' features
- Define squared error,
 - $e_{ij}^2 = (r_{ij} - p_{ij})^2$,
 - $p_{ij} = \sum_k (u_{ik}, v_{kj})$
- Objective: find 'k', minimizing 'e'
- Can achieve using **Gradient Descent**

Implementation

- Obtain gradient by differentiating 'e' w.r.t. 'u' and 'v'
- Update rule,
 - $u_n = u_{n-1} + 2 * a * e_{n-1} * v_{n-1}$
 - $v_n = v_{n-1} + 2 * a * e_{n-1} * u_{n-1}$
- 'a' is the learning rate

Implementation

- **Regularization**
- Modify squared error 'e' as,
 - $e_{ij}^2 = (r_{ij} - p_{ij})^2 + b * \text{sum}_k(|U|_F + |V|_F)$
 - 'b' is regularizer
- Regularized update rule,
 - $u_n = u_{n-1} + 2 * a * (e_{n-1} * v_{n-1} - b * u_{n-1})$
 - $v_n = v_{n-1} + 2 * a * (e_{n-1} * u_{n-1} - b * v_{n-1})$

Pseudo Code

- Initialize U and V randomly
- Until convergence
- For $i = \{1, 2 \dots n\}$
 - For $j = \{1, 2, \dots m\}$
 - $E_{ij} = R_{ij} - \sum_k (U_{ik} V_{kj}) + b * |U|_F * |V|_F$
 - If (error < threshold) **return**
 - $u_{ik} = u_{i,k-1} + 2 * a * (e_{ij} * v_{k-1,j} - b * u_{i,k-1})$
 - $v_{kj} = v_{k-1,j} + 2 * a * (e_{ij} * u_{i,k-1} - b * v_{k-1,j})$

Post-processing

- Bias removal:
 - Correct prediction by adding
 - Correction = $\text{mean}_{\text{training}} - \text{mean}_{\text{prediction}}$
- Prediction Truncation:
 - (Prediction < 1) = 1
 - (Prediction > 5) = 5
- Round-off prediction to nearest integer if difference < 0.1

Results




Fig: Train Error Surface

Results

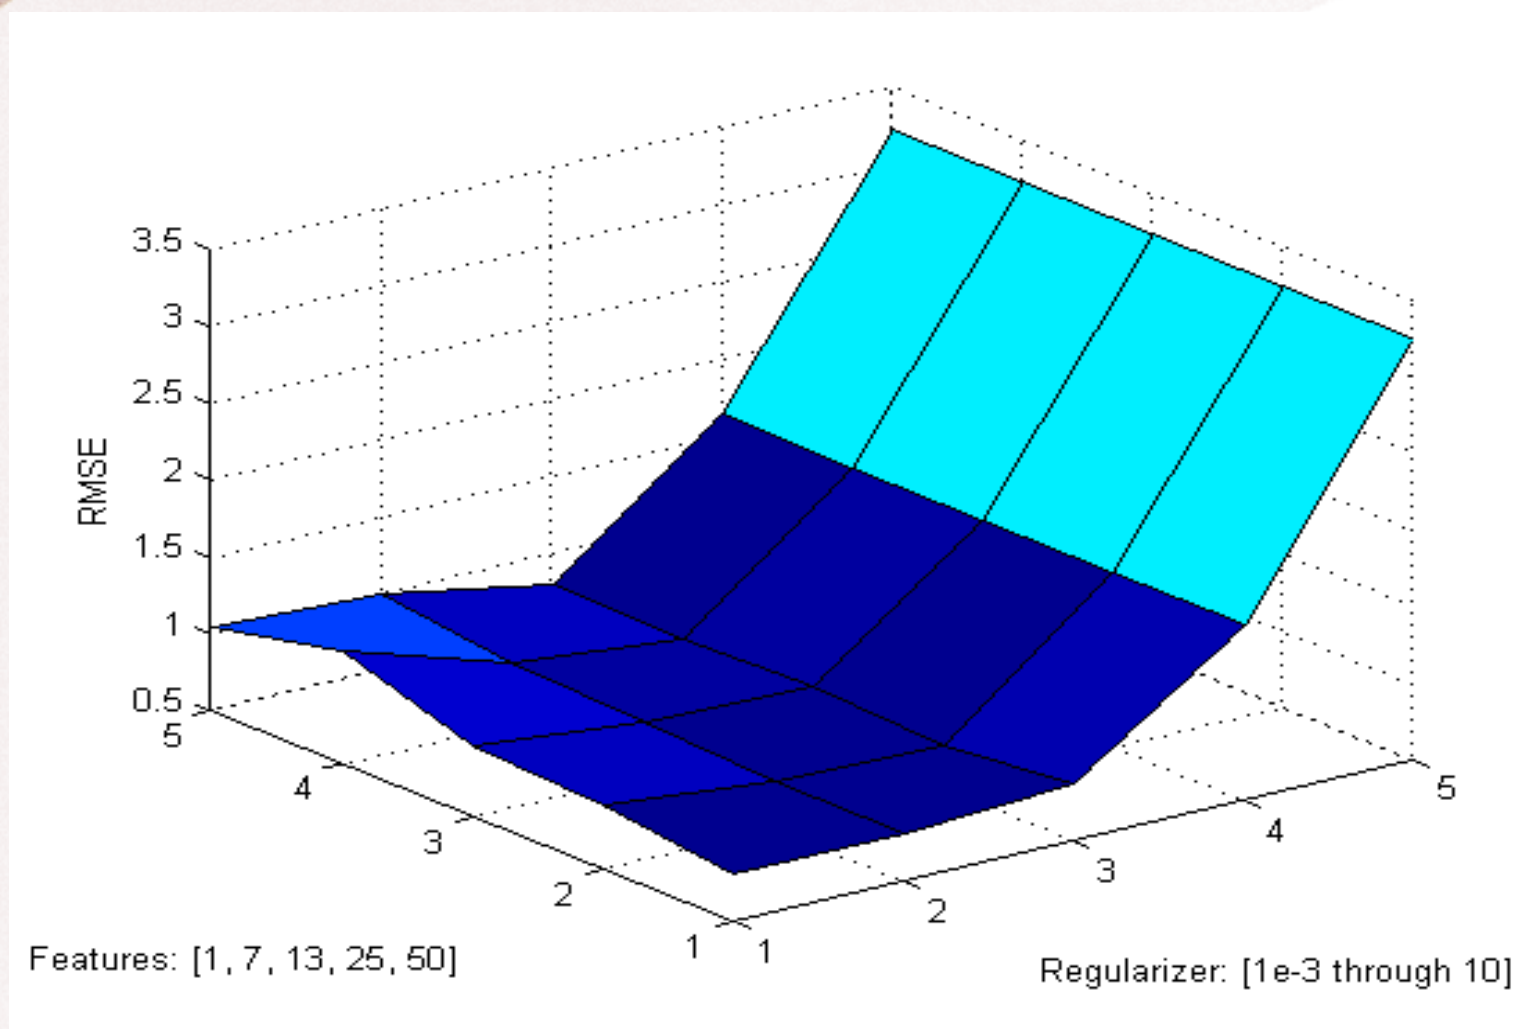
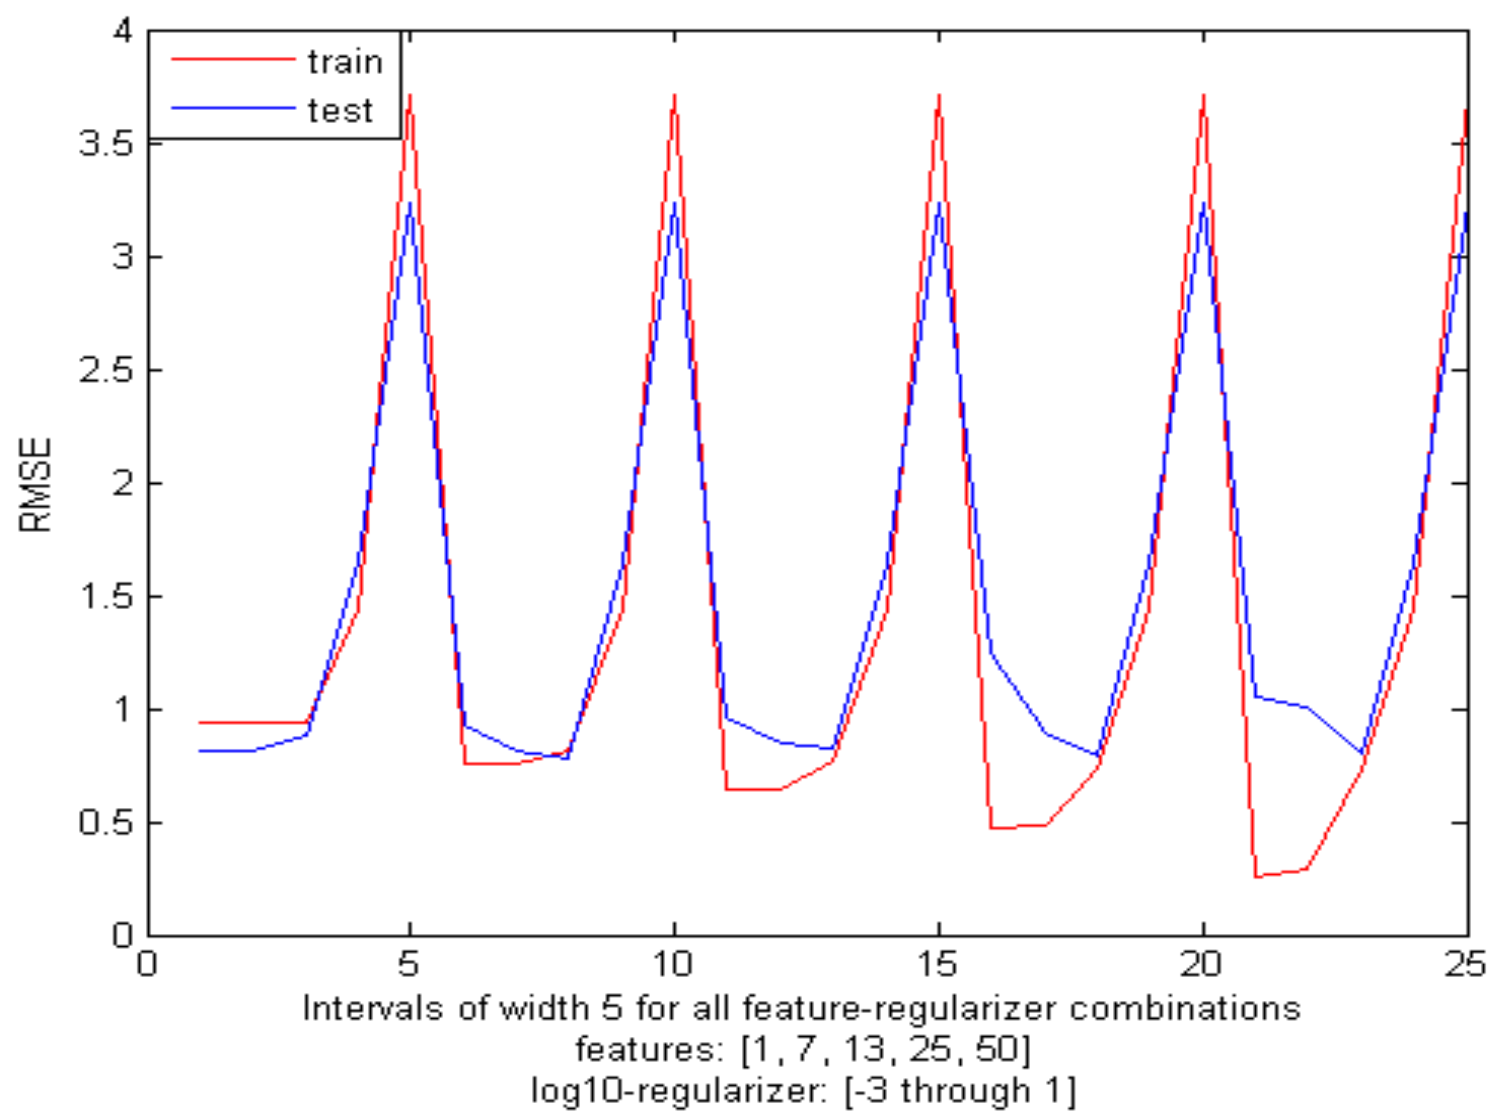



Fig: Test Error Surface

Results



Results




- Features : 175

- Regularizer : 0.04


Experiments

- Algorithm mixture:
 - Compute prediction as convex combination of two algorithms
 - $P = \lambda * P_1 + (1 - \lambda) * P_2$
 - P_1 predicted by 1st algorithm : ALS
 - P_2 predicted by 2nd algorithm : K-NN


Results



$\lambda = 0.05$



$\lambda = 0.5$



$\lambda = 0.95$

References

- “Large-Scale Parallel Collaborative Filtering for the Netflix Prize”, Y. Zhou et.al.
- “Recommendation System Based on Collaborative Filtering”, Zheng Wen